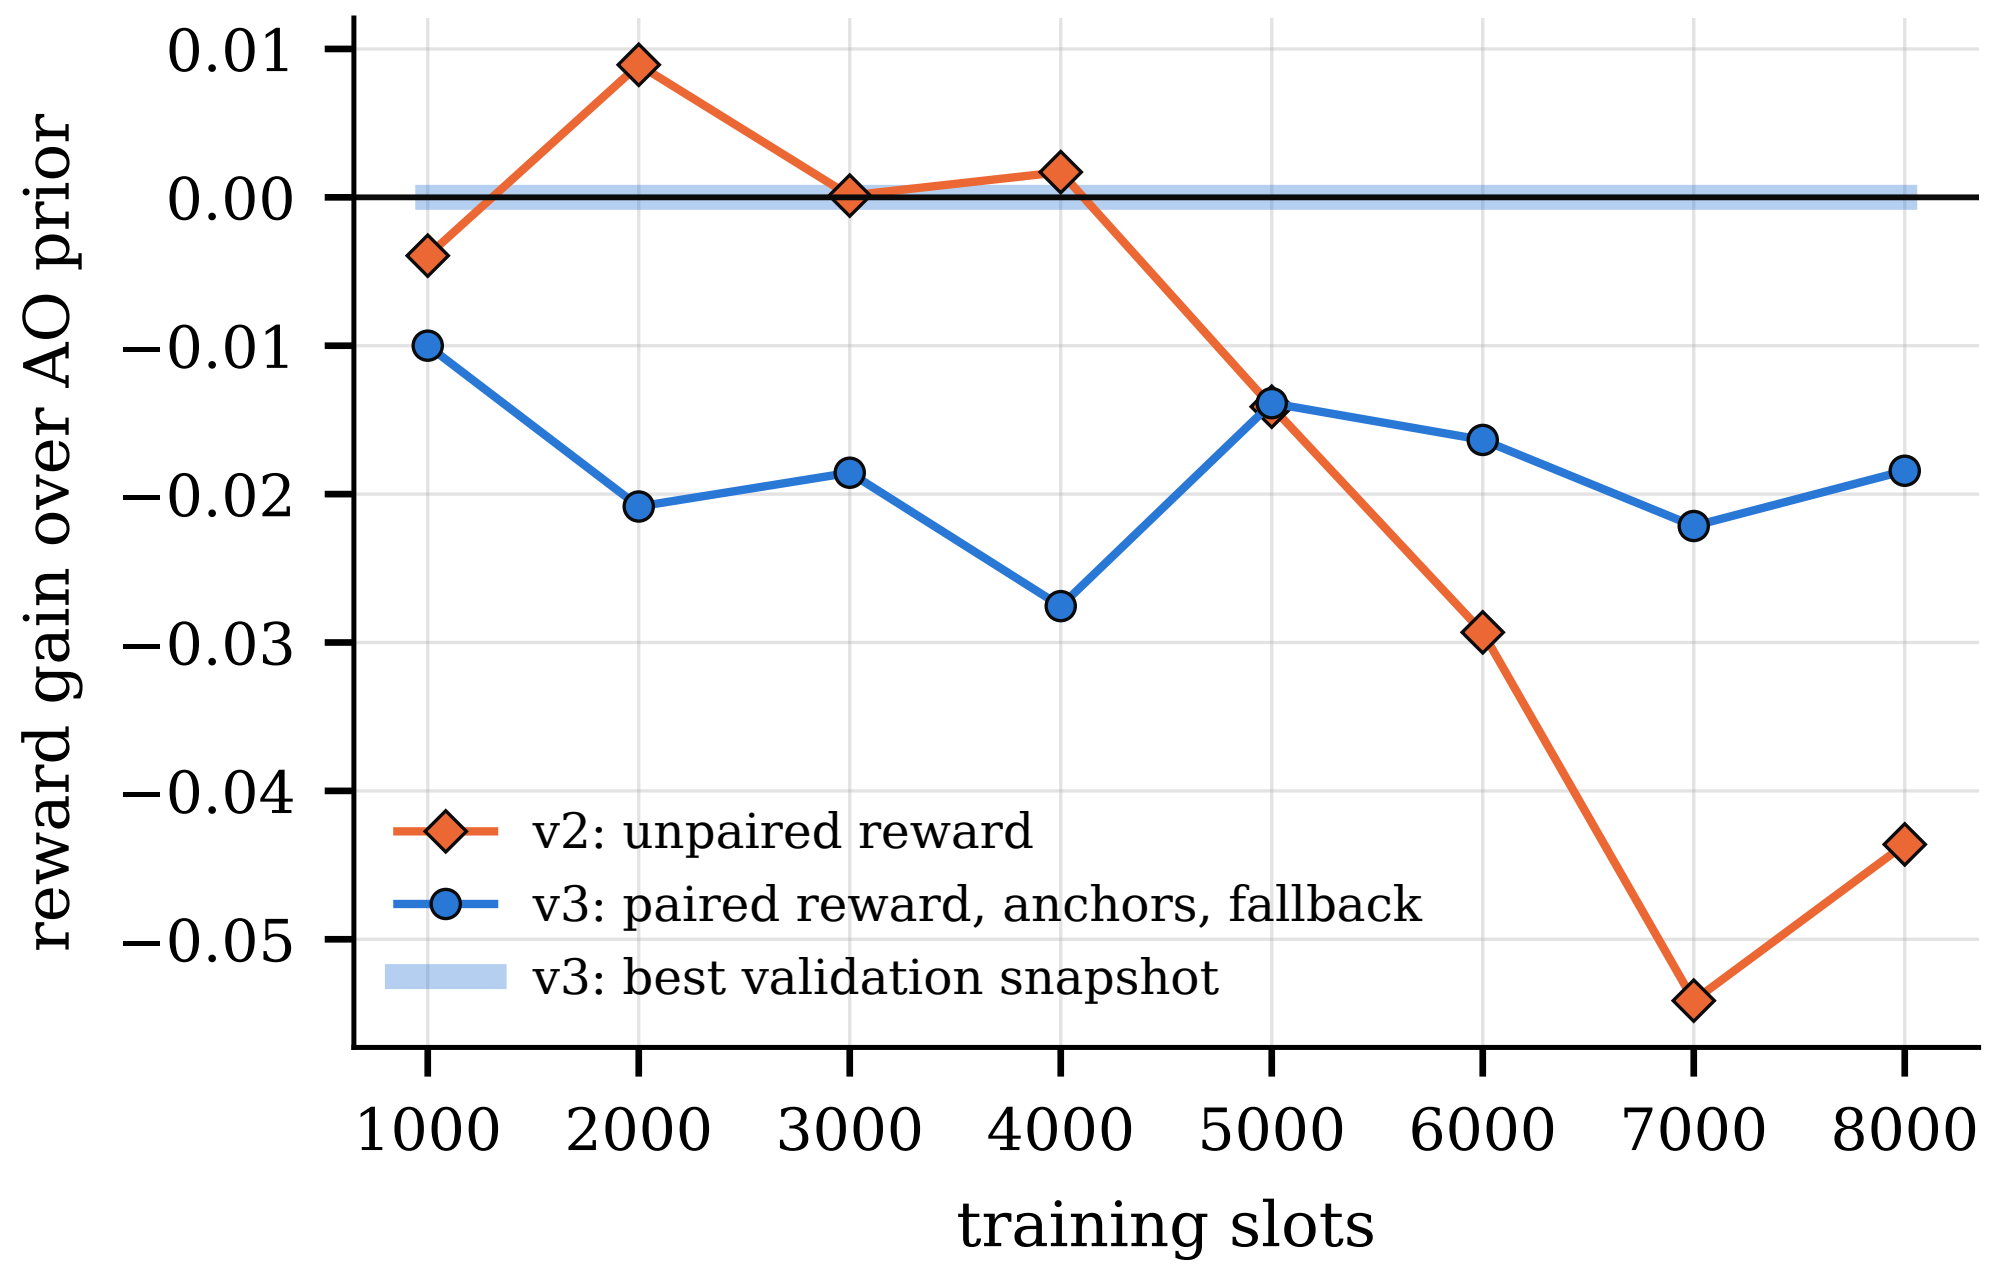
0.01
0.00
−0.01
−0.02
−0.03
−0.04
−0.05
reward gain over AO prior
v2: unpaired reward
v3: paired reward, anchors, fallback
v3: best validation snapshot
1000
2000
3000
4000
5000
6000
7000
8000
training slots